

## Chapter 2 (cont)

### Displaying Data from Multiple Tables

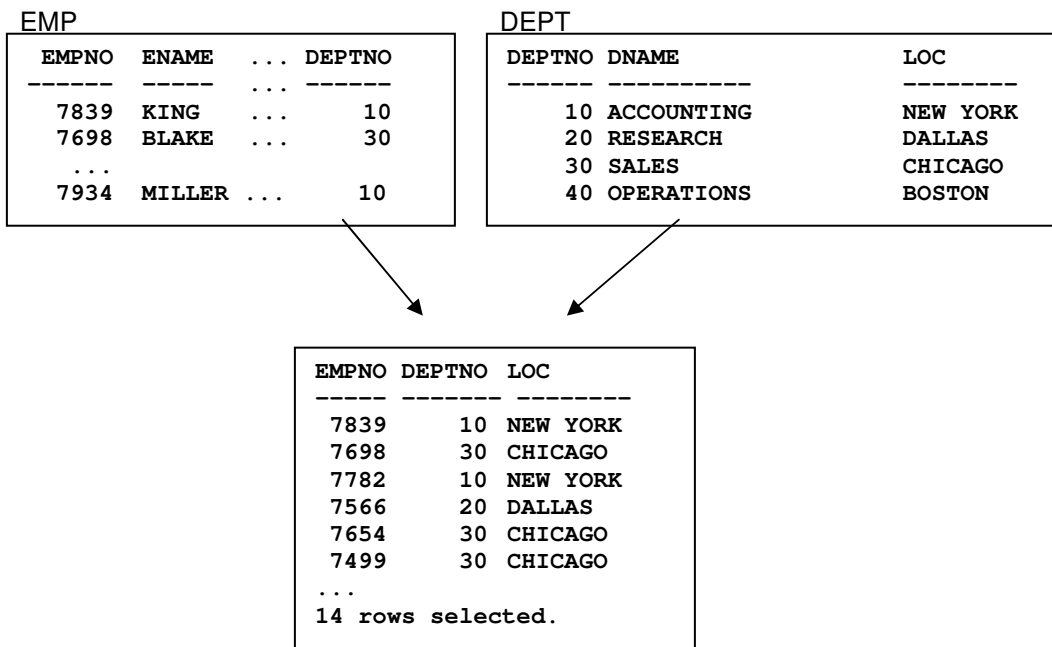
#### Objectives

- Write SELECT statements to access data from more than one table using equality and nonequality joins
- Generate a Cartesian product of all rows from two or more tables
- View data that generally does not meet a join condition by using OUTER joins
- Join a table to itself using a self –join

#### Join Standards

- Before the Oracle 9i release, the join syntax was different from the ANSI standards. This syntax is still very common within the BC Government.
- With Oracle 9i and later releases, can use either the Oracle Join Syntax or the SQL 1999 compliant syntax.
- We will look at both

#### Obtaining Data from Multiple Tables



## What Is a Join?

- Use a join to query data from more than one table.

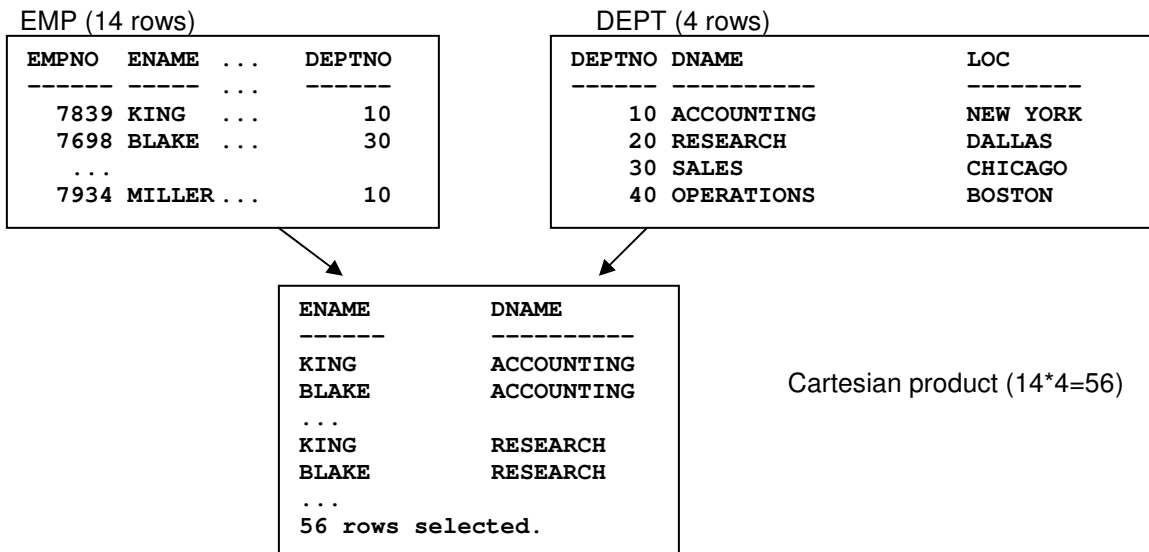
```
SELECT    table1.column, table2.column
FROM      table1, table2
WHERE     table1.column1 = table2.column2;
```

- Write the join condition in the WHERE clause. (Oracle Join Syntax)
- Prefix the column name with the table name when the same column name appears in more than one table.

## Cartesian Product

- A Cartesian product is formed when:
  - A join condition is omitted
  - A join condition is invalid
  - All rows in the first table are joined to all rows in the second table
- To avoid a Cartesian product, always include a valid join condition in a WHERE clause.

## Generating a Cartesian Product



## Cross Joins (99)

- A Cartesian product is also known as a cross join.
- In SQL:1999 use the words CROSS JOIN explicitly to create a cartesian product.
- Note that the join occurs in the FROM clause.

```
SELECT ename, dname
FROM emp CROSS JOIN dept;
```

## Types of Joins

- Equijoin
- Non-equijoin
- Outer join
- Self join

## What Is an Equijoin?

- If the query is relating two tables using an equality operator (=), it is an Equijoin.

EMP

EMPNO	ENAME	DEPTNO
7839	KING	10
7698	BLAKE	30
7782	CLARK	10
7566	JONES	20

↑  
Foreign Key

DEPT

DEPTNO	DNAME	LOC
10	ACCOUNTING	NEW YORK
30	SALES	CHICAGO
10	ACCOUNTING	NEW YORK
20	RESEARCH	DALLAS

↑  
Primary Key

## Retrieving Records with Equijoins


```
SQL> SELECT emp.empno, emp.ename, emp.deptno,  
2      dept.deptno, dept.loc  
3 FROM emp, dept  
4 WHERE emp.deptno=dept.deptno;
```

EMPNO	ENAME	DEPTNO	DEPTNO	LOC
7839	KING	10	10	NEW YORK
7698	BLAKE	30	30	CHICAGO
7782	CLARK	10	10	NEW YORK
7566	JONES	20	20	DALLAS

...  
14 rows selected.

## Qualifying Ambiguous Column Names

- Use table prefixes to qualify column names that are in multiple tables.
- Improve performance by using table prefixes.
- Distinguish columns that have identical names but reside in different tables by using column aliases.



## Using Table Aliases

- Simplify queries by using table aliases.

```
SQL> SELECT emp.empno, emp.ename, emp.deptno,
2         dept.deptno, dept.loc
3 FROM   emp, dept
4 WHERE  emp.deptno=dept.deptno;
```

```
SQL> SELECT e.empno, e.ename, e.deptno,
2         d.deptno, d.loc
3 FROM   emp e, dept d
4 WHERE  e.deptno=d.deptno;
```

## Inner Join

- The ANSI/ISO SQL:1999 standard defines the join we just saw as a Natural Join.
- A inner join is a join based on all the columns in two tables that have the same name and same data types.
- A inner join selects rows from the two tables that have equal values in all the matched columns.

```
SELECT department_name, city
FROM departments INNER JOIN locations;
```

```
SELECT emp.empno, emp.ename, emp.deptno, dept.deptno, dept.loc
FROM   emp, dept
WHERE  emp.deptno=dept.deptno;
```

Becomes:

```
SELECT empno, ename, deptno, loc
FROM   emp INNER JOIN dept;
```

### **The USING clause (99)**

- The USING clause specifies the columns that have to be used for an equijoin between two tables.
- The column name should be the same in both tables, and should have compatible data types.
- The columns reference by the USING clause should not have qualifiers anywhere in the SQL statement, including a WHERE clause.
- For example:
  - You want the last names of all employees and the names of the departments in which they work.
  - Will use the departments and employees tables.

### **The USING clause (99)**

- Can you use a natural join?
- Check the structure of both tables.
- The two tables have two columns in common, manager\_id and department\_id. But only want to join on department\_id.
- Would do the following:

```
SELECT e.last_name, d.department_name
FROM hr.employees e INNER JOIN hr.departments d
USING (department_id);
```
- How would you write this in the previous Oracle way?

### **The ON Clause (99)**

- Can use the ON clause to specify the join condition for the join of two tables.
- This clause separates the join condition from other filter conditions.
- Example:
  - Suppose you need to list the last names of all employees and the departments in which they work.
  - Need to perform an equijoin of the employees and department tables based on the department\_id column.

```
SELECT e.last_name, d.department_name
FROM hr.employees e INNER JOIN hr.departments d
ON (e.department_id = d.department_id);
```
- What is the difference between the USING clause and the ON clause?
  - Can use USING only if the column names are the same in both tables.
  - If the column names are different, must use ON.

## Joining More Than Two Tables

CUSTOMER		ORDER		ITEM	
NAME	CUSTID	CUSTID	ORDID	ORDID	ITEMID
-----	-----	-----	-----	-----	-----
JOCKSPORTS	100	101	610	610	3
TKB SPORT SHOP	101	102	611	611	1
VOLLYRITE	102	104	612	612	1
JUST TENNIS	103	106	601	601	1
K+T SPORTS	105	102	602	602	1
SHAPE UP	106	106	604	602	1
WOMENS SPORTS	107	106	605	...	
...	...	...		64 rows	
9 rows selected.		21 rows selected.		selected.	

...where customer.custid = ord.custid and ord.ordid = item.ordid

### Multitable Join

- Want to join more than two tables.
- In SQL:1999 compliant syntax, join are performed from left to right.
- Example:
  - List the names of employees, the names of the departments in which they work, and the name of the city in which the department is located.
    - Need the employees, departments, and locations table.

```
SELECT employee_id, city, department_name
FROM hr.locations l INNER JOIN hr.departments d
ON (d.location_id = l.location_id)
INNER JOIN hr.employees e
ON (d.department_id = e.department_id);
```

or

```
SELECT employee_id, city, department_name
FROM hr.locations l INNER JOIN hr.departments d
USING (location_id)
INNER JOIN hr.employees e
USING (department_id);
```

- It is important to note that the first join to be performed is *locations INNER JOIN departments*
- This join can reference columns in locations and departments, but not employees.
- The second join can reference columns from all three tables.

## Multitable Join - Oracle

```
SELECT employee_id, city, department_name
FROM hr.locations l, hr.departments d, hr.employees e
WHERE d.location_id = l.location_id
AND d.department_id = e.department_id;
```

## Non-Equijoins

EMP

EMPNO	ENAME	SAL
7839	KING	5000
7698	BLAKE	2850
7782	CLARK	2450
7566	JONES	2975
7654	MARTIN	1250
7499	ALLEN	1600
7844	TURNER	1500
7900	JAMES	950

Want Salaries between HISAL & LOSAL

...  
14 rows selected.

SALGRADE		
GRADE	LOSAL	HISAL
1	700	1200
2	1201	1400
3	1401	2000
4	2001	3000
5	3001	9999

## Retrieving Records with Non-Equijoins

```
SQL> SELECT      e.ename, e.sal, s.grade
2 FROM          emp e, salgrade s
3 WHERE         e.sal
4 BETWEEN      s.losal AND s.hisal;
```

ENAME	SAL	GRADE
JAMES	950	1
MARTIN	1250	2
ALLEN	1600	3
TURNER	1500	3
BLAKE	2850	4 ...

## Outer Joins

- To see data from one table even if there is no corresponding row in the joining table.

No employee in the Operations department →

ENAME	DEPTNO
-----	-----
KING	10
BLAKE	30
CLARK	10
JONES	20
...	

DEPTNO	DNAME
-----	-----
10	ACCOUNTING
30	SALES
10	ACCOUNTING
20	RESEARCH
...	
40	OPERATIONS

## Outer Joins

- You use an outer join to also see rows that do not usually meet the join condition.
- Outer join operator is the plus sign (+).

```
SELECT table1.column, table2.column
FROM   table1, table2
WHERE  table1.column(+) = table2.column;
```

```
SELECT table1.column, table2.column
FROM   table1, table2
WHERE  table1.column = table2.column(+);
```

## Using Outer Joins

```
SQL> SELECT      e.ename, d.deptno, d.dname
2 FROM          emp e, dept d
3 WHERE         e.deptno(+) = d.deptno
4 ORDER BY     e.deptno;
```

ENAME	DEPTNO	DNAME
-----	-----	-----
KING	10	ACCOUNTING
CLARK	10	ACCOUNTING
...		
	40	OPERATIONS

15 rows selected.

Get all departments even if there isn't a matching employee.



## Using Outer Joins

- What would happen for the following?

```
SQL> SELECT          e.ename, d.deptno, d.dname
  2  FROM emp e, dept d
  3  WHERE            e.deptno = d.deptno(+)
  4  ORDER BY        e.deptno;
```

### Outer Join (99)

- In SQL:1999, the syntax of the outer joins has changed.
- Need to know if you want a:
  - Left Outer Join
  - Right Outer Join
  - Full Outer Join

### Left Outer Join (99)

- A left outer join returns all the rows from the table specified on the left side of the JOIN keyword, and the rows that satisfy the join condition from the table that is specified on the right of the join keyword.

### Left Outer Join (99) Example

- Suppose you need a list of all employees and the departments they belong to. This should include the employees who have not yet been assigned to any department.

```
SELECT e.last_name,d.department_name
FROM hr.employees e LEFT OUTER JOIN hr.departments d
ON (e.department_id = d.department_id);
```

- And the equivalent previous Oracle SQL?

### Right Outer Join (99)

- A right outer join returns all the rows from the table specified on the right side of the JOIN keyword, and the rows that satisfy the join condition from the table that is specified on the left of the join keyword.

### Right Outer Join (99) Example

Suppose you need a list of all employees and the departments they belong to. This should include the departments to which no employees have yet been assigned

```
SELECT e.last_name,d.department_name
FROM hr.employees e RIGHT OUTER JOIN hr.departments d
ON (e.department_id = d.department_id);
```

- And the Oracle version?

### Full Outer Join (99)

- A full outer join returns all the rows from the tables that satisfy the join condition as well as the results of both the left and the right outer joins.
- This is new to Oracle9i +

### Full Outer Join (99) Example

- Suppose you need a list of all employees and the departments they belong to. This should include the employees who have not yet been assigned to any department.

```
SELECT e.last_name,d.department_name
FROM hr.employees e FULL OUTER JOIN hr.departments d
ON (e.department_id = d.department_id);
```

- No direct previous Oracle version.
- Would need to use a UNION as follows:

```
SELECT e.last_name,d.department_name
FROM hr.employees e, hr.departments d
WHERE e.department_id(+) = d.department_id
UNION
SELECT e.last_name,d.department_name
FROM hr.employees e, hr.departments d
WHERE e.department_id = d.department_id(+);
```

## Self Joins

- Joining a Table to Itself

EMP(WORKER)

EMPNO	ENAME	MGR
-----	-----	----
7839	KING	
7698	BLAKE	7839
7782	CLARK	7839
7566	JONES	7839
7654	MARTIN	7698
7499	ALLEN	7698

EMP(MANAGER)

EMPNO	ENAME
-----	-----
7839	KING
7839	KING
7839	KING
7698	BLAKE
7698	BLAKE

“MGR in the WORKER table is equal to EMPNO in the MANAGER table”

## Joining a table to itself

```
SQL> SELECT worker.ename||' works for '||manager.ename
2   FROM      emp worker, emp manager
3   WHERE      worker.mgr = manager.empno;
```

```
WORKER.ENAME || 'WORKSFOR' || MANAG
-----
BLAKE works for KING
CLARK works for KING
JONES works for KING
MARTIN works for BLAKE
...
13 rows selected.
```

## Joining a Table to Itself (99)

```
SELECT e.last_name || ' works for ' || m.last_name
FROM hr.employees e INNER JOIN hr.employees m
ON (e.manager_id = m.employee_id);
```

## Set Operators

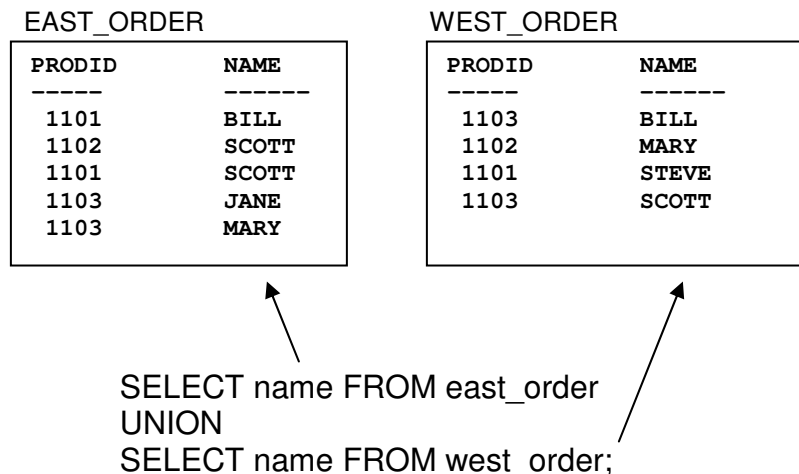
- Used to select data from multiple tables
- Combine the results of two queries into one
- Are compound queries
- All set operators have equal precedence – if multiples, evaluate from left to right
- Data types of the resulting columns should match in both queries

## Set Operators

- Oracle has four set operators
  - UNION: returns all unique rows selected by either query
  - UNION ALL: returns all rows, including duplicates selected by either query
  - INTERSECT: returns rows selected from both queries
  - MINUS: returns unique rows selected by the first query but not the rows selected by the second query

## UNION

- All unique rows selected by either query.



## UNION ALL

- ALL rows selected by either query.

```
SELECT name FROM east_order  
UNION  
SELECT name FROM west_order;
```

## INTERSECT

- Rows selected from both queries

```
SELECT name FROM east_order
INTERSECT
SELECT name FROM west_order;
```

## MINUS

- Unique rows selected by first query but not rows selected by second.

```
SELECT name FROM east_order
MINUS
SELECT name FROM west_order;
```

## SUMMARY

```
SELECT      table1.column, table2.column
FROM        table1, table2
WHERE       table1.column1 = table2.column2;
```

For joins: Equijoin, Non-equijoin, Outer join, Self join

```
SELECT      column FROM      table
{UNION, UNION ALL, INTERSECT, MINUS}
SELECT      column FROM      table;
```

## SET OPERATORS