

## Chapter 8

### Changing Table Contents

#### Objectives

- After completing this lesson, you should be able to do the following:
  - Describe each DML statement
  - Insert rows into a table
  - Update rows in a table
  - Delete rows from a table
  - Control transactions

#### Data Manipulation Language

- A DML statement is executed when you:
  - Add new rows to a table
  - Modify existing rows in a table
  - Remove existing rows from a table
- A *transaction* consists of a collection of DML statements that form a logical unit of work.

#### Adding a New Row to a Table

##### The INSERT Statement

Add new rows to a table by using the INSERT statement.

```
INSERT INTO      table [(column [, column...])]  
VALUES          (value [, value...]);
```

- Only one row is inserted at a time with this syntax.

#### Inserting New Rows

- Insert a new row containing values for each column.
- List values in the default order of the columns in the table.
- Optionally list the columns in the INSERT clause.

```
SQL> INSERT INTO dept  
      2 (deptno, dname, loc)  
      3 VALUES  
      4 (50, 'DEVELOPMENT', 'DETROIT');  
1 row created.
```

- Enclose character and date values within single quotation marks.

#### Inserting Rows with Null Values

- Implicit method: Omit the column from the column list.

```
SQL> INSERT INTO dept  
      2 (deptno, dname )  
      3 VALUES  
      4 (60, 'MIS');  
1 row created.
```

- Explicit method: Specify the NULL keyword.

```
SQL> INSERT INTO dept
2 VALUES
3 (70, 'FINANCE', NULL);
1 row created.
```

### Inserting Special Values

- The SYSDATE function records the current date and time.

```
SQL> INSERT INTO emp
2 (empno, ename, job,
3 mgr, hiredate, sal, comm,
4 deptno)
5 VALUES
6 (7196, 'GREEN', 'SALESMAN',
7 7782, SYSDATE, 2000, NULL,
8 10);
1 row created.
```

### Inserting Specific Date Values

- Add a new employee.

```
SQL> INSERT INTO emp
2 VALUES
3 (2296, 'AROMANO', 'SALESMAN', 7782,
4 TO_DATE('FEB 3, 97', 'MON DD, YY'),
5 1300, NULL, 10);
1 row created.
```

### Inserting Values by Using Substitution Variables

- Create an interactive script by using SQL\*Plus substitution parameters.

```
SQL> INSERT INTO dept
2 (deptno, dname, loc)
3 VALUES
4 (&department_id,
5 '&department_name', '&location');
```

```
Enter value for department_id: 80
Enter value for department_name: EDUCATION
Enter value for location: ATLANTA
```

1 row created.

### Copying Rows from Another Table

- Write your INSERT statement with a subquery.

```
SQL> INSERT INTO managers(id, name, salary, hiredate)
  2      SELECT      empno, ename, sal, hiredate
  3      FROM emp
  4      WHERE      job = 'MANAGER';
3 rows created.
```

- Do not use the VALUES clause.
- Match the number of columns in the INSERT clause to those in the subquery.

### Changing Data in a Table

#### The UPDATE Statement

- Modify existing rows with the UPDATE statement.

```
UPDATE      table
SET          column = value [, column = value, ...]
[WHERE      condition];
```

- Update more than one row at a time, if required.

#### Updating Rows in a Table

- Specific row or rows are modified when you specify the WHERE clause.

```
SQL> UPDATE      emp
  2 SET      deptno = 20
  3 WHERE empno = 7782;
1 row updated.
```

- All rows in the table are modified if you omit the WHERE clause.

```
SQL> UPDATE      employee
  2 SET      deptno = 20;
14 rows updated.
```

#### Updating with Multiple-Column Subquery

- Update employee 7698's job and department to match that of employee 7499.

```
SQL> UPDATE emp
  2 SET  (job, deptno) =
  3      (SELECT job, deptno
  4      FROM emp
  5      WHERE empno = 7499)
  6 WHERE empno = 7698;
1 row updated.
```

### Updating Rows Based on Another Table

- Use subqueries in UPDATE statements to update rows in a table based on values from another table.

```
SQL> UPDATE      employee
  2 SET          deptno = (SELECT deptno
  3              FROM emp
  4              WHERE empno = 7788)
  5 WHERE        job = (SELECT job
  6              FROM emp
  7              WHERE empno = 7788);
2 rows updated.
```

### Updating Rows: Integrity Constraint Error

- Department number 55 does not exist

```
SQL> UPDATE      emp
  2 SET          deptno = 55
  3 WHERE        deptno = 10;
```

```
UPDATE emp
*
```

ERROR at line 1:

ORA-02291: integrity constraint (USR.EMP\_DEPTNO\_FK) violated - parent key not found

### Removing a Row from a Table The DELETE Statement

- You can remove existing rows from a table by using the DELETE statement.

```
DELETE [FROM]      table
[WHERE      condition];
```

### Deleting Rows from a Table

- Specific rows are deleted when you specify the WHERE clause.

```
SQL> DELETE FROM  department
  2 WHERE          dname = 'DEVELOPMENT';
  1 row deleted.
```

- All rows in the table are deleted if you omit the WHERE clause.

```
SQL> DELETE FROM  department;
4 rows deleted.
```

### Deleting Rows Based on Another Table

- Use subqueries in DELETE statements to remove rows from a table based on values from another table.

```
SQL> DELETE FROM  employee
  2 WHERE          deptno =
  3              (SELECT deptno
  4              FROM dept
  5              WHERE dname ='SALES');
6 rows deleted.
```

### Deleting Rows: Integrity Constraint Error

```
SQL> DELETE FROM dept
      2 WHERE deptno = 10;
```

```
DELETE FROM dept
      *
```

ERROR at line 1:

ORA-02292: integrity constraint (USR.EMP\_DEPTNO\_FK) violated - child record found

—You cannot delete a row that contains a primary key that is used as a foreign key in another table.

### Truncate Statement

- If just want to empty a table of all rows, can use the DDL statement TRUNCATE.
- Is like a DELETE statement without a WHERE clause.
- Will remove ALL rows from a table.
- Unlike DELETE, cannot rollback.
  - DDL statements will implicitly perform a commit. This includes all pending DML statements.

### Truncate vs Delete

- TRUNCATE is very fast as does not generate undo information.
- When a table is truncated, the storage for the table and all indexes can be reset back to the initial size. DELETE will never shrink the size of a table or its indexes.

### Using Default Values

- Book says that this was introduced in Oracle9i, but has been around longer than that. Definitely in 8i.
- Can define the default value as part of the CREATE TABLE statement.

### MERGE Statement

- Use the MERGE statement to select rows from one table for update or insertion into another table.
- The decision whether to update or insert into the target table is based on a condition in the ON clause.
- This statement is a convenient way to combine at least two operations. Can avoid multiple INSERT and UPDATE DML statements.
- MERGE is a deterministic statement. That is, you cannot update the same row of the target table multiple times in the same MERGE statement.
- Is available in Oracle & SQL Server etc but not MySQL as yet.

- Syntax:

```
MERGE [hint] INTO [schema .] table [t_alias] USING [schema .]
{ table | view | subquery } [t_alias] ON ( condition )
  WHEN MATCHED THEN merge_update_clause
  WHEN NOT MATCHED THEN merge_insert_clause;
```

- Assume have a table with the following:

SELECT \* FROM bonuses;

EMPLOYEE\_ID BONUS

```
-----
153          100
154          100
155          100
156          100
158          100
159          100
160          100
161          100
163          100
```

- *The Human Resources manager decides that all employees should receive a bonus. Those who have not made sales get a bonus of 1% of their salary. Those who already made sales get an increase in their bonus equal to 1% of their salary. The MERGE statement implements these changes in one step:*

*MERGE INTO bonuses D*

```
  USING (SELECT employee_id, salary, department_id FROM employees WHERE
  department_id = 80) S
  ON (D.employee_id = S.employee_id)
  WHEN MATCHED THEN UPDATE SET D.bonus = D.bonus + S.salary*.01
  WHEN NOT MATCHED THEN INSERT (D.employee_id, D.bonus) VALUES
  (S.employee_id, S.salary*0.1);
```

SELECT \* FROM bonuses;

EMPLOYEE\_ID BONUS

EMPLOYEE_ID	BONUS
153	180
154	175
155	170
156	200
158	190
159	180
160	175
161	170
163	195
157	950
145	1400
170	960
179	620
152	900
169	1000

### Database Transactions

- Consist of one of the following statements:
  - DML statements that make up one consistent change to the data
  - One DDL statement
  - One DCL statement
- Begin when the first executable SQL statement is executed
- End with one of the following events:
  - COMMIT or ROLLBACK is issued
  - DDL or DCL statement executes (automatic commit)
  - User exits
  - System crashes

### Advantages of COMMIT and ROLLBACK Statements

- Ensure data consistency
- Preview data changes before making changes permanent
- Group logically related operations

### Implicit Transaction Processing

- An automatic commit occurs under the following circumstances:
  - DDL statement is issued
  - DCL statement is issued
  - Normal exit from SQL\*Plus, without explicitly issuing COMMIT or ROLLBACK
- An automatic rollback occurs under an abnormal termination of SQL\*Plus or a system failure.

### State of the Data Before COMMIT or ROLLBACK

- The previous state of the data can be recovered.
- The current user can review the results of the DML operations by using the SELECT statement.
- Other users *cannot* view the results of the DML statements by the current user.
- The affected rows are *locked*; other users cannot change the data within the affected rows.

### State of the Data After COMMIT

- Data changes are made permanent in the database.
- The previous state of the data is permanently lost.
- All users can view the results.
- Locks on the affected rows are released; those rows are available for other users to manipulate.

### Committing Data

- Make the changes.

```
SQL> UPDATE emp
2   SET      deptno = 10
3   WHERE    empno = 7782;
1 row updated.
```

- Commit the changes.

```
SQL> COMMIT;
Commit complete.
```

### State of the Data After ROLLBACK

- Discard all pending changes by using the ROLLBACK statement.
  - Data changes are undone.
  - Previous state of the data is restored.
  - Locks on the affected rows are released.

```
SQL> DELETE FROM employee;
14 rows deleted.
```

```
SQL> ROLLBACK;
Rollback complete.
```



## Rolling Back Changes to a Marker

- Create a marker in a current transaction by using the SAVEPOINT statement.
- Roll back to that marker by using the ROLLBACK TO SAVEPOINT statement.

```
SQL> UPDATE...
```

```
SQL> SAVEPOINT update_done;
```

Savepoint created.

```
SQL> INSERT...
```

```
SQL> ROLLBACK TO update_done;
```

Rollback complete.

## Statement-Level Rollback

- If a single DML statement fails during execution, only that statement is rolled back.
- The Oracle Server implements an implicit savepoint.
- All other changes are retained.
- The user should terminate transactions explicitly by executing a COMMIT or ROLLBACK statement.

## Read Consistency

- Read consistency guarantees a consistent view of the data at all times.
- Changes made by one user do not conflict with changes made by another user.
- Read consistency ensures that on the same data:
  - Readers do not wait for writers
  - Writers do not wait for readers

## Locking

- Oracle locks:
  - Prevent destructive interaction between concurrent transactions
  - Require no user action
  - Automatically use the lowest level of restrictiveness
  - Are held for the duration of the transaction
  - Have two basic modes:
    - Exclusive
    - Share

## Summary

Statement	Description
INSERT	Adds a new row to the table
UPDATE	Modifies existing rows in the table
DELETE	Removes existing rows from the table
COMMIT	Makes all pending changes permanent
SAVEPOINT	Allows a rollback to the savepoint marker
ROLLBACK	Discards all pending data changes